

The pthread library

- In Linux, when a new process is created, it already contains a thread, used to execute the main() function. Additional threads can be created using the pthread library, which is part of the C library.
- Of course all threads inside a given process will share the same address space, the same set of open files, etc.
- The pthread library also provides thread synchronization primitives: mutexes and conditions.
- This pthread library has its own header: pthread.h
- Applications using pthread function calls should be explicitly linked with the pthread library: gcc o app app.c lpthread

Creating a new thread:

The function to create a new thread is pthread_create()

```
int pthread_create(pthread_t * thread, pthread_attr_t * attr, void *(*start_routine)(void *), void * arg);
```

- thread is a pointer to a pthread_t structure that will be initialized by the function. Later, this structure can be used to reference the thread.
- Attr is a pointer to an optional structure pthread_attr_t. This structure can be manipulated using pthread_attr_*() functions. It can be used to set various attributes of the threads (detach policy, scheduling policy, etc.)
- start_routine is the function that will be executed by the thread
- arg is the private data passed as argument to the start_routine function.

Thread creation. code sample:

```
#include <pthread.h>
void *thread(void *data)
{
    while(1) {
        printf(« Hello world from thread »);
    }
}

int main(void) {
    pthread_t th;
    pthread_create(& th, NULL, thread, NULL);
    return 0;
}
```

Joinable and detached threads

- When the main() function exits, all threads of the application are destroyed.
- The pthread_join() function call can be used to suspend the execution of a thread until another thread terminates. This function must be called in order to release the resources used by the thread, otherwise it remains as zombie.
- Threads can also be detached, in which case they become independent. This can be achieved using.
 - Thread attributes at thread creation, using:

```
pthread_attr_setdetachstate(& attr, PTHREAD_CREATE_DETACHED);
```
 - pthread_detach(), passing the pthread_t structure as argument.

Thread join. code sample

```
#include <pthread.h>
void *thread(void *data)
{
    int i;
    for (i = 0; i < 100; i++) {
        printf(« Hello world from thread »);
    }
}

int main(void) {
    pthread_t th;
    pthread_create(& th, NULL, thread, NULL);
    pthread_join(& th, NULL);
    return 0;
}
```

Thread cancelation

- It is also possible to cancel a thread from another thread using the `pthread_cancel()` function, passing the `pthread_t` structure of the thread to cancel.

```
#include <pthread.h>
```

```
void *thread(void *data)
```

```
{  
    while(1) {  
        printf(« Hello world from thread »);  
    }  
}
```

```
int main(void) {  
    pthread_t th;  
    pthread_create(& th, NULL, thread, NULL);  
    sleep(1);  
    pthread_cancel(& th);  
    pthread_join(& th, NULL);  
    return 0;  
}
```

pthread mutexes

- The `pthread` library provides a mutual exclusion primitive, the `pthread_mutex`.

- Declaration and initialization of a `pthread` mutex.

Solution 1: at definition time: `pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;`

Solution 2: at runtime: `pthread_mutex_t lock;`

```
...  
pthread_mutex_init(& lock, NULL);  
...
```

```
pthread_mutex_destroy(& lock);
```

- The second argument to `pthread_mutex_init()` is a set of mutex-specific attributes, in the form of a `pthread_mutexattr_t` structure that can be initialized and manipulated using `pthread_mutexattr_*` functions.

- Take the mutex: `ret = pthread_mutex_lock(& lock);`

- If the mutex is already taken by the calling threads, three possible behaviours depending on the mutex type (defined at creation time)

- **Normal (« fast ») mutex:** the function doesn't return, deadlock
- **« Error checking » mutex:** the function return with the `EDEADLK` error
- **« Recursive mutex »:** the function returns with success

- Release the mutex: `ret = pthread_mutex_unlock(& lock);`

- Try to take the mutex: `ret = pthread_mutex_trylock(& lock);`

pthread conditions

- Conditions can be used to suspend a thread until a condition becomes true, as signaled by another thread.

- Initialization, static or dynamic

- `pthread_cond_t cond = PTHREAD_COND_INITIALIZER;`
- `pthread_cond_t cond;`
- `pthread_cond_init(& cond, NULL);`

- Wait for the condition: `pthread_cond_wait(& cond, & mutex);` The mutex will be released before waiting and taken again after the wait.

- Signaling the condition

- To one thread waiting: `pthread_cond_signal(& cond);`
- To all threads waiting: `pthread_cond_broadcast(& cond);`

pthread conditions example:

Receiver side:

```
pthread_mutex_lock(& lock);
```

```
while(is_queue_empty())
```

```
    pthread_cond_wait(& cond, & lock);
```

```
/* Something in the queue, and we have the mutex ! */
```

```
pthread_mutex_unlock(& lock);
```

Sender side:

```
pthread_mutex_lock(& lock);  
/* Add something to the queue */  
pthread_mutex_unlock(& lock);  
pthread_cond_signal(& cond);
```

POSIX shared memory

- A great way to communicate between processes without going through expensive system calls.
- Open a shared memory object:
 - `shm_fd = shm_open("acme", O_CREAT | O_RDWR, 0666);`
A zero size `/dev/shm/acme` file appears.
- Set the shared memory object size:
 - `ftruncate(shm_fd, SHM_SIZE);`
`/dev/shm/acme` is now listed with the specified size.
- If the object has already been sized by another process, you can get its size with the `fstat` function.
- Map the shared memory in process address space:
 - `addr = mmap (0, SHM_SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);`
Now we have a memory area we can use!
- Lock the shared memory in RAM (best for realtime tasks):
 - `mlock(addr, SHM_SIZE);`
- Use the shared memory object! Other processes can use it too.

Exiting:

- Unmap the shared memory object:
 - `munmap (addr, SHM_SIZE);`
This automatically unlocks it too.
- Close it:
 - `close (shm_fd);`
- Remove the shared memory object:
 - `shm_unlink ("acme");`
The object is effectively deleted after the last call to `shm_unlink`.

POSIX message queues:

Deterministic and efficient IPC. See `man mqueue.h`. Advantages for realtime applications:

- Preallocated message buffers
- Messages with priority. A message with a higher priority is always received first.
- Send and receive functions are synchronous by default. Possibility to set a wait timeout to avoid nondeterminism.
- Support asynchronous delivery notifications.

Creating and opening a message queue

- Declare queue attributes:
 - `queue_attr.mq_maxmsg = 16; /* max number of messages in queue */`
 - `queue_attr.mq_msgsize = 128; /* max message size */`
- Open a queue:
 - `qd = mq_open("/msg_queue", O_CREAT | O_RDWR, 0600, &queue_attr);`
- Posting a message:
 - `#define PRIORITY 3`
`char msg[] = "Goodbye Bill";`
`mqsend(qd, msg, strlen(msg), PRIORITY);`
- Closing the queue:
 - `mq_close(qd);`

Receiving a message:

- Opening the shared message queue:
 - `qd = mq_open("/msg_queue", O_RDWR, 0600, NULL);`
- Waiting for a message:
 - `mq_receive(qd, text, buf, buf_size, &prio);`
- Close the queue:
 - `mq_close(qd);`
- Destroy the queue:
 - `mq_unlink("/msg_queue");`

POSIX semaphores:

Resources for sharing resources between threads or processes. See `man semaphore.h`.

- Named semaphores: can be used between unrelated processes.
- Unnamed semaphores: can be used between threads from the same process, or by related processes (parent / child).
- Open and / or create a named semaphore:
 - `sem_open();`
- Close a named semaphore:
 - `sem_close();`
- Destroy a named semaphore:
 - `sem_unlink();`
- Initialize an unnamed semaphore:
 - `sem_init();`
- Destroy an unnamed semaphore:
 - `sem_destroy();`
- Get current semaphore count:
 - `sem_getvalue();`
- Try to lock the semaphore. Wait otherwise:
 - `sem_wait();`
- Just tries to lock the semaphore, but gives up if the semaphore is already locked:
 - `sem_trywait();`
- Release the semaphore:
 - `sem_post();`

POSIX signals:

- Signals are a mechanism to notify a process that an event occurred : expiration of a timer, completion of an asynchronous I/O operation, or any kind of event specific to your application
- Signals are also used internally by the system to tell a process that it must be suspended, restarted, stopped, that is has done an invalid memory reference, etc.
- Each signal is identified by a number : `SIGSEGV`, `SIGKILL`, `SIGUSR1`, etc.
- An API is available to catch signals, wait for signals, mask signals, etc.

Registering a signal handler

- A signal handler can be registered using: `sighandler_t signal(int signum, sighandler_t handler);`
⇒ The handler has the following prototype : `void handler(int signum)`
- `int sigaction(int signum, const struct sigaction *act, struct sigaction *oldact);`
⇒ The `sigaction` structure contains the reference to the handler
⇒ The handler can have two different prototypes:
 - `void handler(int signum)`
 - `void handler(int signum, siginfo_t *info, void *data)`
- Inside the handler code, only some functions can be used : only the `asyncsignalsafe` functions.

Signal registration example:

```
#include <signal.h>
#include <assert.h>
#include <unistd.h>
#include <stdio.h>
```

```
void myhandler(int signum){

    printf("Signal caught!\n");

}
```

```
int main(void){
    int ret;
    struct sigaction action = {
        .sa_handler = myhandler,
    };
    ret = sigaction(SIGUSR1, & action, NULL);
```

```
assert(ret == 0);
while(1);
return 0;
```

```
}
```

⇒ From the command line, the signal can then be sent using `kill USR1 PID`

Sending a signal:

- From the command line, with the famous `kill` command, specifying the `PID` of the process to which the signal should be sent. By default, `kill` will send `SIGTERM`. Another signal can be sent using `kill USR1`.

- POSIX provides a function to send a signal to a process: `int kill(pid_t pid, int sig);`

⇒ In a multithread program, the signal will be delivered to an arbitrary thread. Use `tkill()` to send the signal to a specific thread.

Signal sets and their usage

- A type `sigset_t` is defined by POSIX, to hold a set of signals. This type is manipulated through different functions:

- `sigemptyset()` to empty the set of signals
- `sigaddset()` to add a signal to a set
- `sigdelset()` to remove a signal from a set
- `sigfillset()` to fill the set of signals with all signals

- Signals can then be blocked or unblocked using: `sigprocmask(int how, const sigset_t *set, sigset_t *oldset);`

- `sigset_t` are also used in many other functions:

- `sigaction()` to give the list of signals that must be blocked during execution of the handler.
- `sigpending()` to get the list of pending signals

Waiting for signals:

2 ways of waiting for signals:

- `sigwaitinfo()` and `sigtimedwait()` to wait for blocked signals (signals which remain pending until they are processed by a thread waiting for them.)

- `sigsuspend()` to register a signal handler and suspend the thread until the delivery of an unblocked signal (which are delivered without waiting for a thread to wait for them).

POSIX clocks and timers

Compared to standard (BSD) timers in Linux

- Possibility to have more than 1 timer per process.
- Increased precision, up to nanosecond accuracy
- Timer expiration can be notified either with a signal or with a thread.
- Several clocks available.

Available POSIX clocks:

Defined in `/usr/include/linux/time.h`

- `CLOCK_REALTIME`: Systemwide clock measuring the time in seconds and nanoseconds since Jan 1, 1970, 00:00.

Can be modified. Accuracy: 1/HZ (1 to 10 ms)

- `CLOCK_MONOTONIC`: Systemwide clock measuring the time in seconds and nanoseconds since system boot.

Cannot be modified, so can be used for accurate time measurement. Accuracy: 1/HZ

- `CLOCK_PROCESS_CPUTIME_ID`: Measures process uptime. 1/HZ accuracy. Can be changed.

- `CLOCK_THREAD_CPUTIME_ID`: Same, but only for the current thread.

Time management:

Functions defined in `time.h`

- `clock_settime`: Set the specified clock to a value
- `clock_gettime`: Read the value of a given clock
- `clock_getres`: Get the resolution of a given clock.

Using timers

Functions also defined in `time.h`

- `clock_nanosleep`: Suspend the current thread for the specified time, using a specified clock.
- `nanosleep`: Same as `clock_nanosleep`, using the `CLOCK_REALTIME` clock.
- `timer_create`: Create a timer based on a given clock.
- `timer_delete`: Delete a timer
- `timer_settime`: Arm a timer.
- `timer_gettime`: Access the current value of a timer.

Asynchronous I/O:

- Helpful to implement nonblocking I/O. Allows to overlap compute tasks with I/O processing, to increase determinism.

- Supported functionality:

- Send multiple I/O requests at once from different sources
- Cancel ongoing I/O requests
- Wait for request completion
- Inquire the status of a request: completed, failed, or in progress.

Compiling instructions:

- Includes: nothing special to do. Available in the standard path.
- Libraries: link with librt. Example: gcc lrt o rttest rttest.c